

PANDAS

(in depth - 1)

Data Manipulation

Can be considered of comprising 3 stages:

- ❑ Data preparation
- ❑ Data transformation
- ❑ Data aggregation

This lecture will look at Data Preparation!

Data preparation

- ❑ **Merge()**: connects the rows in a dataframe based on one or more keys (somewhat like Join in SQL).
- ❑ **Concat()**: concatenates the objects along an axis.
- ❑ **combine_first()**: connects overlapping data in a dataframe to enable the filling in of missing values by taking data from another structure.
- ❑ **Pivot**: enables the interchange of rows and columns in a dataframe.

Setting up

```
>>> import numpy as np
>>> import pandas as pd
```

```
>>> frame1 = pd.DataFrame( {'id':['ball','pencil','pen','mug','ashtray'], 'price':
                             [12.33,11.44,33.21,13.23,33.62]})
```

```
>>> frame2 = pd.DataFrame( {'id':['pencil','pencil','ball','pen'], 'color': ['white','red','red','black']})
```

```
>>> print(frame1)
```

	id	price
0	ball	12.33
1	pencil	11.44
2	pen	33.21
3	mug	13.23
4	ashtray	33.62

```
>>> print(frame2)
```

	color	id
0	white	pencil
1	red	pencil
2	red	ball
3	black	pen

Merge operation

To merge the 2 dataframes:

```
>>>> merged_frame = pd.merge(frame1,frame2)
```

the resulting DataFrame consists of all rows that **have an ID in common** between the two DataFrames.

```
>>> print(merged_frame)
```

	id	price	color
0	ball	12.33	red
1	pencil	11.44	white
2	pencil	11.44	red
3	pen	33.21	black

Usually you need to specify the merge column:

```
>>> pd.merge(frame1,frame2,on='id')
```

```
>>> frame1 = pd.DataFrame(
{'id':['ball','pencil','pen','mug','ashtray'],
 'color': ['white','red','red','black','green'],
 'brand': ['OMG','ABC','ABC','POD','POD']})
```

```
>>> frame1
```

	brand	color	id
0	OMG	white	ball
1	ABC	red	pencil
2	ABC	red	pen
3	POD	black	mug
4	POD	green	ashtray

```
>>> frame2 = pd.DataFrame(
{'id':['pencil','pencil','ball','pen'],
 'brand': ['OMG','POD','ABC','POD']})
```

```
>>> frame2
```

	brand	id
0	OMG	pencil
1	POD	pencil
2	ABC	ball
3	POD	pen

```
>>> pd.merge(frame1,frame2)
```

Empty DataFrame

Columns: [brand, color, id]

Index: []

Since both columns of frame2 are present in frame1. Ambiguity !

Use 'on' option to explicitly define the criterion of merging that pandas must follow

```
>>> pd.merge(frame1,frame2,on='id')
```

	brand_x	color	id	brand_y
0	OMG	white	ball	ABC
1	ABC	red	pencil	OMG
2	ABC	red	pencil	POD
3	ABC	red	pen	POD

```
>>> pd.merge(frame1,frame2,on='brand')
```

	brand	color	id_x	id_y
0	OMG	white	ball	pencil
1	ABC	red	pencil	ball
2	ABC	red	pen	ball
3	POD	black	mug	pencil
4	POD	black	mug	pen
5	POD	green	ashtray	pencil
6	POD	green	ashtray	pen

What if key columns in two DataFrames do not have the same name?

- ❑ use the **left_on** and **right_on** options that specify the key column for the first and for the second DataFrame.

```
>>> frame1 = pd.DataFrame(  
    {'id':['ball','pencil','pen','mug','ashtray'],  
    'color': ['white','red','red','black','green'],  
    'brand': ['OMG','ABC','ABC','POD','POD']})
```

```
>>> frame1
```

	brand	color	id
0	OMG	white	ball
1	ABC	red	pencil
2	ABC	red	pen
3	POD	black	mug
4	POD	green	ashtray


```
>>> frame2 = pd.DataFrame(  
    {'sid':['pencil','pencil','ball','pen'],  
    'brand': ['OMG','POD','ABC','POD']})
```

```
>>> frame2
```

	brand	sid
0	OMG	pencil
1	POD	pencil
2	ABC	ball
3	POD	pen

```
>>> pd.merge(frame1, frame2, left_on='id', right_on='sid')
```

	brand_x	color	id	brand_y	sid
0	OMG	white	ball	ABC	ball
1	ABC	red	pencil	OMG	pencil
2	ABC	red	pencil	POD	pencil
3	ABC	red	pen	POD	pen



Here, id is the join column for the first dataframe and sid the join column for the second dataframe.

The sql equivalent would be: `where frame1.id = frame2.sid`

- ✓ By default, the `merge()` function performs an inner join; the keys in the result are the result of an intersection.

Merge contd..

- ❑ Other possible merge operations are the left join, the right join, and the outer join.
- ❑ The outer join produces the union of all keys, combining the effect of a left join with a right join.
- ❑ To select the type of join you have to use the “how” option.

```
>>> pd.merge(frame1,frame2,on='id',how='outer')    # ensures all rows included from both frame even if they don't match
```

- ❑ Driving the merge from the left:

```
>>> pd.merge(frame1,frame2,on='id',how='left')    # all rows from frame1 and any rows from frame2 that match
```

- ❑ Or from the right:

```
>>> pd.merge(frame1,frame2,on='id',how='right')    # all rows from frame2 and any rows from frame1 that match
```

To merge multiple keys, simply add a list to the on option:

```
>>> pd.merge(frame1,frame2,on=['id','brand'],how='outer')
```

See the tutorial for more examples !

Concatenating

- NumPy has a concatenate function for concatenating arrays:

```
>>> array1 = np.array([[0, 1, 2], [3, 4, 5], [6, 7, 8]])
```

```
>>> array2 = np.arange(9).reshape((3,3))+6
```

```
>>> array3 = np.concatenate([array1,array2],axis=1)
```

```
print(array3)
```

```
array([[ 0,  1,  2,  6,  7,  8],  
       [ 3,  4,  5,  9, 10, 11],  
       [ 6,  7,  8, 12, 13, 14]])
```

Concatenating rows of arrays

```
>>> np.concatenate([array1,array2],axis=0)
```

```
array([[ 0,  1,  2],  
       [ 3,  4,  5],  
       [ 6,  7,  8],  
       [ 6,  7,  8],  
       [ 9, 10, 11],  
       [12, 13, 14]])
```

The Pandas concat() function

```
>>> ser1 = pd.Series(np.random.rand(4), index=[1,2,3,4])  
>>> ser2 = pd.Series(np.random.rand(4), index=[5,6,7,8])
```

```
print(ser1)
```

```
1    0.636584  
2    0.345030  
3    0.157537  
4    0.070351  
dtype: float64
```

```
print(ser2)
```

```
5    0.411319  
6    0.359946  
7    0.987651  
8    0.329173  
dtype: float64
```

```
>>> ser3 = pd.concat([ser1,ser2])
```

```
print(ser3)
```

```
1    0.636584  
2    0.345030  
3    0.157537  
4    0.070351  
5    0.411319  
6    0.359946  
7    0.987651  
8    0.329173  
dtype: float64
```

The Pandas concat() function..

- ❑ By default, the concat() function works on axis = 0, returning a series object.
- ❑ If you set the axis = 1, then the result will be a DataFrame.

```
>>> ser3 = pd.concat([ser1,ser2],axis=1)
```

```
print(ser3)
```

	0	1
1	0.636584	NaN
2	0.345030	NaN
3	0.157537	NaN
4	0.070351	NaN
5	NaN	0.411319
6	NaN	0.359946
7	NaN	0.987651
8	NaN	0.329173

This has performed an outer join. This can be changed by setting the join **option to 'inner'**:

```
>>> pd.concat([ser1,ser2],axis=1,join='inner')
```

Empty dataframe

Columns: [0,1]

Index = []

What will be output for

```
>>> pd.concat([ser1,ser3],axis=1,join='inner') ??
```

The Pandas concat() function..

- ❑ To create a hierarchical index on the axis of concatenation we need to use the keys option:

```
>>> pd.concat([ser1,ser2], keys=[1,2])
```

```
1      1    0.636584
      2    0.345030
      3    0.157537
      4    0.070351
2      5    0.411319
      6    0.359946
      7    0.987651
      8    0.329173
```

```
dtype: float64
```

Concatenating dataframes

- ❑ In the case of combinations between Series along the axis = 1, the keys become the column headers of the DataFrame.
- ❑ Essentially applies the same approach as we saw with series objects:

```
>>> frame1 = pd.DataFrame(np.random.rand(9).reshape(3,3), index=[1,2,3], columns=['A','B','C'])  
>>> frame2 = pd.DataFrame(np.random.rand(9).reshape(3,3), index=[4,5,6], columns=['A','B','C'])
```

```
>>> pd.concat([frame1, frame2])
```

	A	B	C
1	0.400663	0.937932	0.938035
2	0.202442	0.001500	0.231215
3	0.940898	0.045196	0.723390
4	0.568636	0.477043	0.913326
5	0.598378	0.315435	0.311443
6	0.619859	0.198060	0.647902

defaults to rows

Concatenating dataframes ..

For columns:

```
>>> pd.concat([frame1, frame2], axis=1)
```

	A	B	C	A	B	C
1	0.400663	0.937932	0.938035	NaN	NaN	NaN
2	0.202442	0.001500	0.231215	NaN	NaN	NaN
3	0.940898	0.045196	0.723390	NaN	NaN	NaN
4	NaN	NaN	NaN	0.568636	0.477043	0.913326
5	NaN	NaN	NaN	0.598378	0.315435	0.311443
6	NaN	NaN	NaN	0.619859	0.198060	0.647902

Combine()

If we wish the two datasets to have indexes that overlap in their entirety or at least partially, we can use `combine_first()`.

Lets define two series:

```
>>> ser1 = pd.Series(np.random.rand(5),index=[1,2,3,4,5])  
>>> ser2 = pd.Series(np.random.rand(4),index=[2,4,5,6])
```

```
print(ser1)
```

```
1    0.942631  
2    0.033523  
3    0.886323  
4    0.809757  
5    0.800295  
dtype: float64
```

```
print(ser2)
```

```
2    0.739982  
4    0.225647  
5    0.709576  
6    0.214882  
dtype: float64
```

```
>>> ser1.combine_first(ser2)
```

```
1    0.942631  
2    0.033523  
3    0.886323  
4    0.809757  
5    0.800295  
6    0.214882  
dtype: float64
```

Combine() ..

- ❑ If you want a partial overlap, you can specify only the portion of the Series you want to overlap.

```
>>> ser1[:3].combine_first(ser2[:3])
```

```
1    0.942631
```

```
2    0.033523
```

```
3    0.886323
```

```
4    0.225647
```

```
5    0.709576
```

```
dtype: float64
```

Pivoting with Hierarchical Indexing

In the context of pivoting there are two basic operations:

- ❑ **Stacking**: rotates or pivots the data structure converting **columns to rows**
- ❑ **Unstacking**: converts **rows into columns**

```
>>> frame1 = pd.DataFrame(np.arange(9).reshape(3,3), index=['white','black','red'],  
                           columns=['ball','pen','pencil'])
```

```
>>> print(frame1)
```

	ball	pen	pencil
white	0	1	2
black	3	4	5
red	6	7	8

Pivoting with hierarchical indexing ..

- Using the `stack()` function on the DataFrame, pivots the columns into rows, thus producing a series:

```
>>> ser5 = frame1.stack()
>>> print(ser5)
```

white	ball	0
	pen	1
	pencil	2
black	ball	3
	pen	4
	pencil	5
red	ball	6
	pen	7
	pencil	8

dtype: int32

- From this [hierarchically indexed series](#), you can reassemble the DataFrame into a pivoted table by use of the `unstack()` function.

```
>>> ser5.unstack()
```

	ball	pen	pencil
white	0	1	2
black	3	4	5
red	6	7	8

Pivoting with hierarchical indexing ..

- ❑ You can also do the unstack on a different level, specifying the number of levels or its name as the argument of the function.

```
>>> ser5.unstack(0)
```

	white	black	red
ball	0	3	6
pen	1	4	7
pencil	2	5	8

Removing columns and rows

Lets define a dataframe:

```
>>> frame1 = pd.DataFrame(np.arange(9).reshape(3,3), index=['white','black','red'],  
                           columns=['ball','pen','pencil'])
```

```
>>> frame1
```

	ball	pen	pencil
white	0	1	2
black	3	4	5
red	6	7	8

❑ To **remove a column**, simply use the del command applied to the DataFrame with the column name specified

```
>>> del frame1['ball']           # removes column "ball" from frame1
```

❑ To **remove an unwanted row**, you have to use the drop() function with the label of the corresponding index as argument

```
>>> frame1.drop('white')        # removes the first row "white" from frame1
```